

*In The Name Of God*

## Subject: Generative Models

Presenter: Reza Karimzadeh


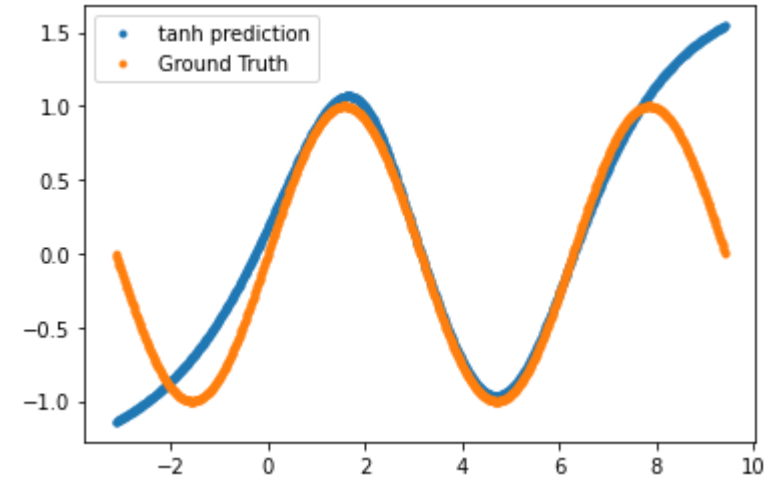


اولین کنگره بین المللی  
فناوریهای پیشرفته در حوزه سلامت و کاربردهای هوش مصنوعی در پزشکی  
The First International Congress on  
Advanced Health Technologies-Artificial Intelligence in Medicine





# Supervised Learning

- ❖ Available pairs of data (X, Y)
  - X: input
  - Y: label
- ❖ Learn a function for mapping X to Y
- ❖ Examples:
  - regression
  - classification
  - segmentation
  - image captioning
  - Etc.



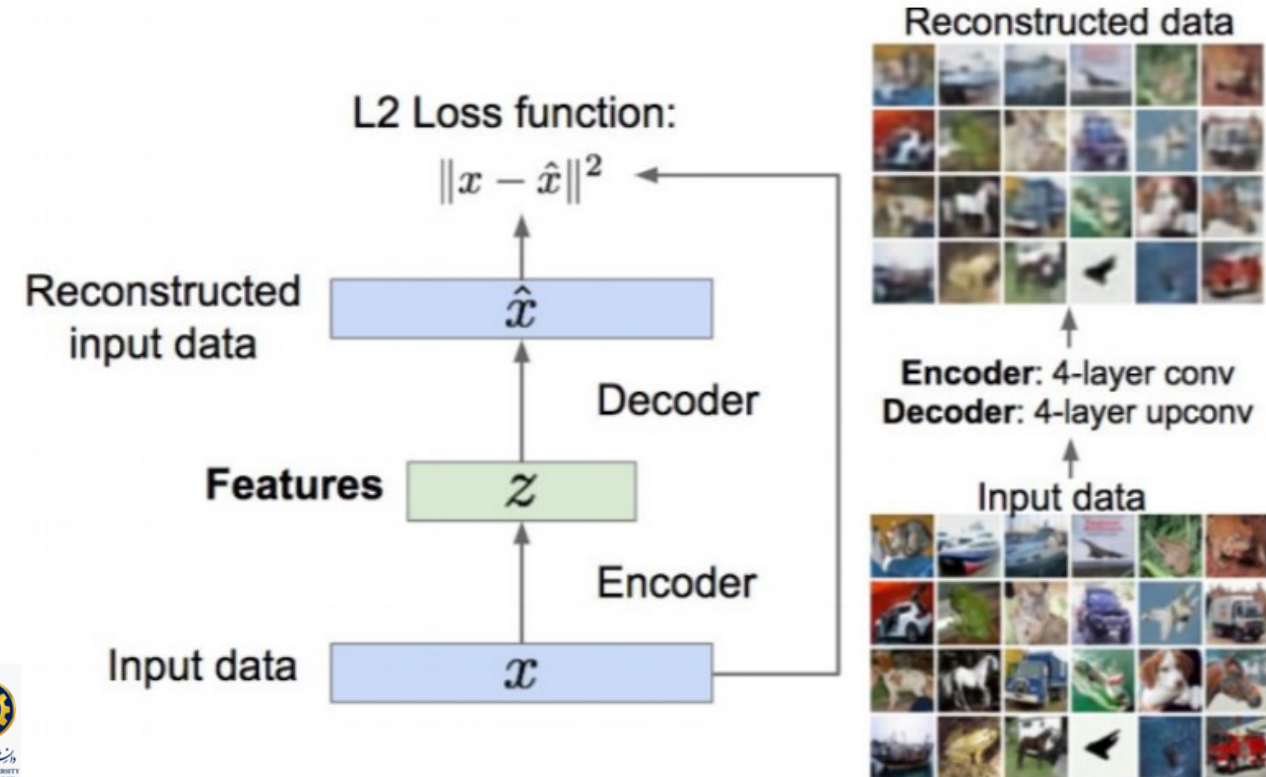
# Unsupervised Learning

- ❖ Available data just X
  - No label
- ❖ Learn hidden structure of data
- ❖ Examples:
  - clustering
  - dimensionality reduction
  - feature learning
  - etc.



# Autoencoders

- ❖ map input data ( $X$ ) to a lower dimension ( $Z$ )
- ❖ reconstruct data ( $\hat{X}$ ) from extracted features ( $Z$ )
- ❖ L2-norm for cost function



# Autoencoders


❖ consider a simple AE

$$\hat{X} = W^T W X$$

$$Loss = \|\hat{X} - X\|^2 = \|W^T W X - X\|^2$$

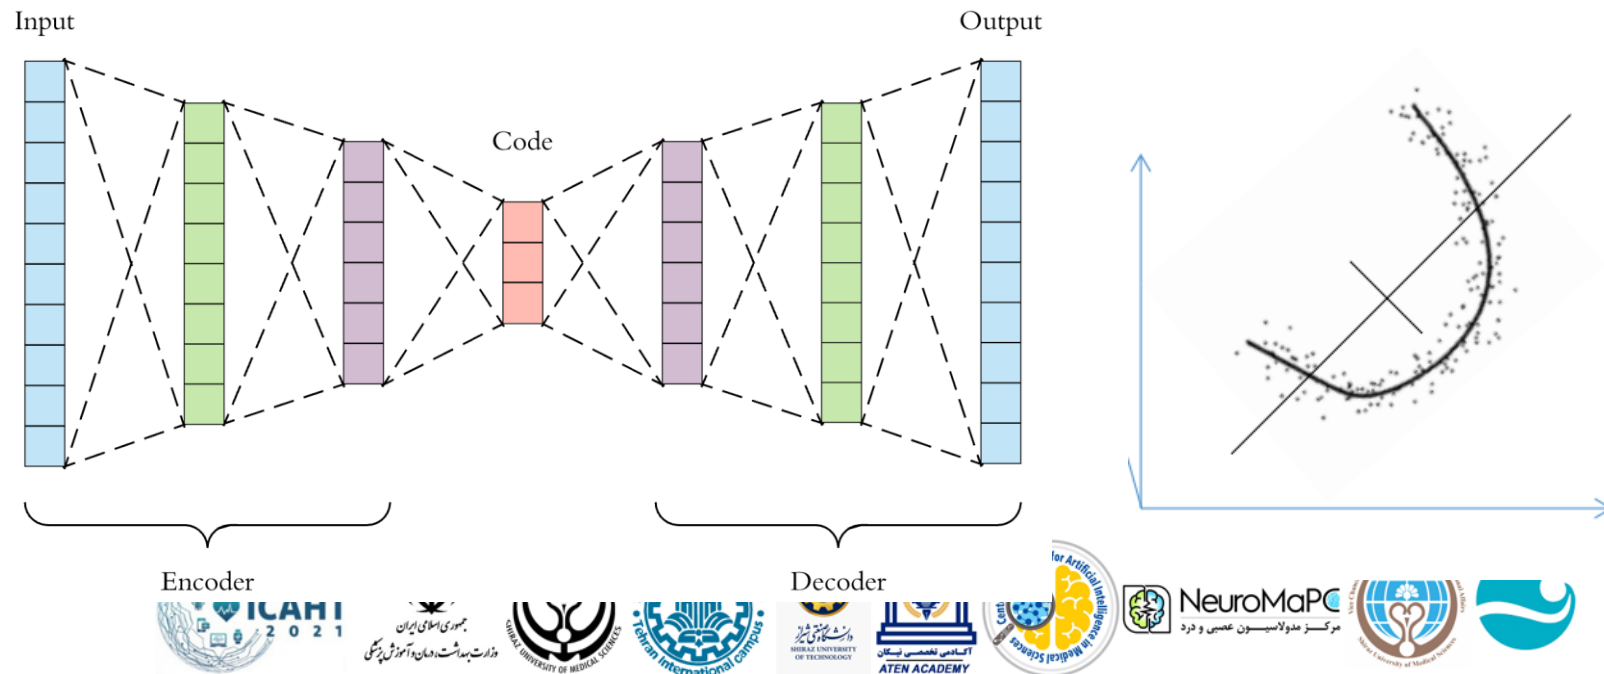
$$W^* = \min_W \|W^T W X - X\|^2$$

❖ It is PCA !!



# Autoencoders



- ❖ We add some nonlinearity to Autoencoders i.e. activation function
  - Nonlinear PCA
- ❖ Deeper networks can capture more complicated manifolds → deep autoencoder





# Autoencoders


- ❖ AE learn to reconstruct sin function in  $[0, 2\pi]$
- ❖ if we keep decoder part and vary the latent space (Z) the output will be a sin function in  $[0, 2\pi]$  but in other parts is non-sin
- ❖ Is it a generative model?



# Autoencoders

## ❖ Is it a generative model?

- Decoder only can generate data that lie on training data manifold
- Samples generated from distribution of training data
- Samples are similar to training data






# Autoencoders (applications)

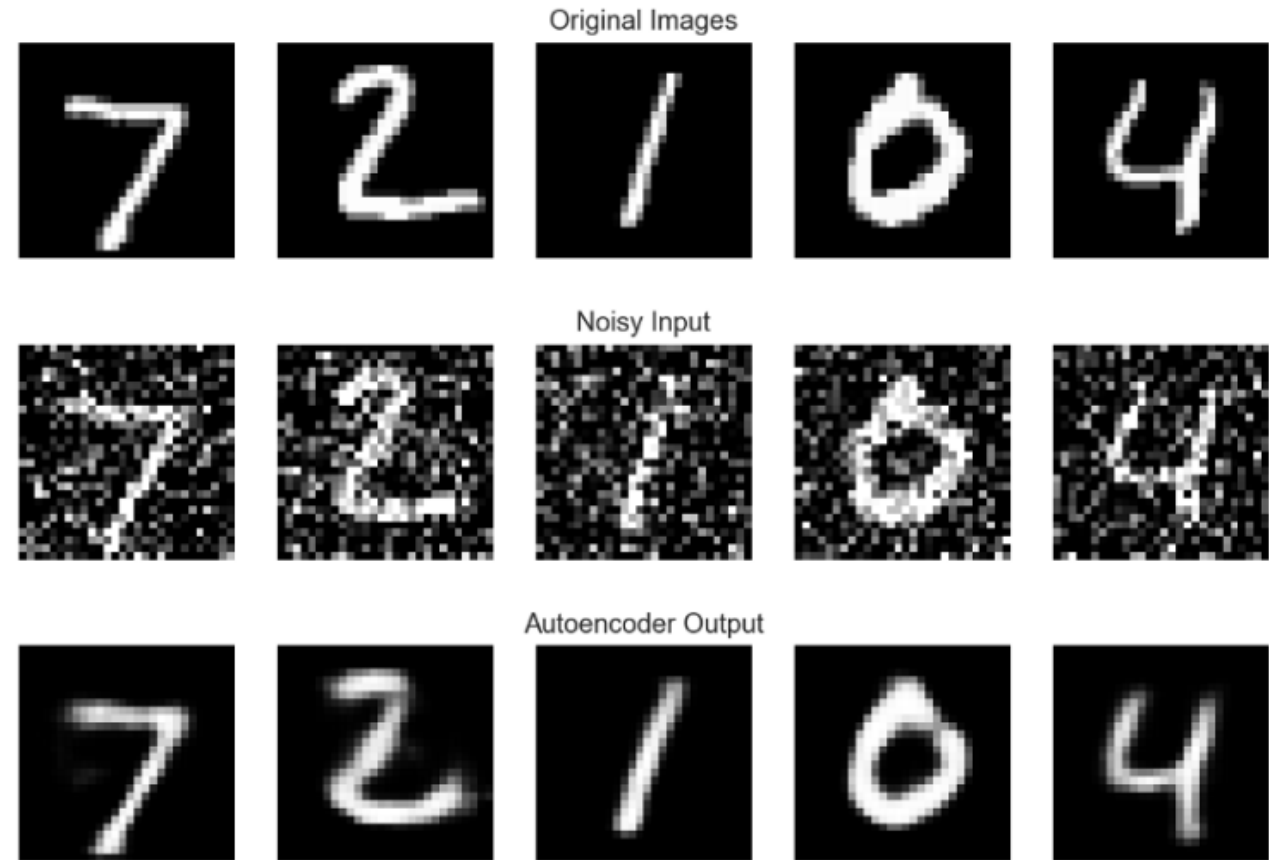
❖ Low dimensional feature extraction

784D  $\rightarrow$  32D  $\rightarrow$  784D




# Autoencoders (applications)

## ❖ Denoising




# Autoencoders (applications)

## ❖ Image reconstruction



# Autoencoders (applications)

## ❖ Super Resolution



Coupled Deep Autoencoder for Single Image Super-Resolution, IEEE, ITC, 2015



# Variational Autoencoder



# Generative Models

- ❖ Explicitly working with the distribution
  - Variational autoencoders
- ❖ Implicit generative model (Generative Adversarial Networks)
  - trained without even needing to explicitly define a density functions



# Variational Autoencoder

Let's define some notions:

1.  $X$ : data that we want to model a.k.a the animal
2.  $z$ : latent variable a.k.a our imagination
3.  $P(X)$ : probability distribution of the data, i.e. that animal kingdom
4.  $P(z)$ : probability distribution of latent variable, i.e. our brain, the source of our imagination
5.  $P(X|z)$ : distribution of generating data given latent variable, e.g. turning imagination into real animal





# Variational Autoencoder

- ❖ Our objective here is to model the data, hence we want to find  $P(X)$ . Using the law of probability, we could find it in relation with  $z$  as follows:

$$P(X) = \int P(X|z)P(z)dz$$

- ❖ VAE idea: Infer  $P(z)$  from  $P(z|X)$ 
  - estimate latent variable likely using our data
- ❖ How to infer  $P(z|X)$  ?
  - infer  $P(z|X)$  using a method called Variational Inference (VI)
  - The main idea of VI is to pose the inference by approach it as an optimization problem By modeling the true distribution  $P(z|X)$  using simpler distribution that is easy to evaluate, e.g. Gaussian,



# Variational Autoencoder

❖ Let infer  $P(z|X)$  using  $Q(z|X)$  via KL distance:

$$\begin{aligned} D_{KL}[Q(z|X) || P(z|X)] &= \sum_z Q(z|X) \log \frac{Q(z|X)}{P(z|X)} \\ &= E \left[ \log \frac{Q(z|X)}{P(z|X)} \right] \\ &= E[\log Q(z|X) - \log P(z|X)] \end{aligned}$$



# Variational Autoencoder

❖ Let use  $P(X)$ ,  $P(X|z)$ , and  $P(z)$ , via Bayes Rule

$$\begin{aligned} D_{KL}[Q(z|X) || P(z|X)] &= E \left[ \log Q(z|X) - \log \frac{P(X|z)P(z)}{P(X)} \right] \\ &= E[\log Q(z|X) - (\log P(X|z) + \log P(z) - \log P(X))] \\ &= E[\log Q(z|X) - \log P(X|z) - \log P(z) + \log P(X)] \end{aligned}$$

Expectation is over  $z$ , this moves  $P(X)$  outside



# Variational Autoencoder

$$D_{KL}[Q(z|X)||P(z|X)] = E[\log Q(z|X) - \log P(X|z) - \log P(z)] + \log P(X)$$

$$D_{KL}[Q(z|X)||P(z|X)] - \log P(X) = E[\log Q(z|X) - \log P(X|z) - \log P(z)]$$

- ❖ If we look carefully at the right hand side of the equation, we would notice that it could be rewritten as another KL divergence.



# Variational Autoencoder

- ❖ rewritten another KL divergence on the right hand side by rearranging sign.

$$D_{KL}[Q(z|X)||P(z|X)] - \log P(X) = E[\log Q(z|X) - \log P(X|z) - \log P(z)]$$

$$\begin{aligned}\log P(X) - D_{KL}[Q(z|X)||P(z|X)] &= E[\log P(X|z) - (\log Q(z|X) - \log P(z))] \\ &= E[\log P(X|z)] - E[\log Q(z|X) - \log P(z)] \\ &= E[\log P(X|z)] - D_{KL}[Q(z|X)||P(z)]\end{aligned}$$



# Variational Autoencoder

❖ VAE objective function:

$$\log P(X) - D_{KL}[Q(z|X)||P(z|X)] = E[\log P(X|z)] - D_{KL}[Q(z|X)||P(z)]$$

❖ So we have got:

1.  $Q(z|X)$  that project our data  $X$  into latent variable space  $\rightarrow$  **Encoder**
2.  $z$ , the latent variable  $\rightarrow$  **Generated code**
3.  $P(X|z)$  that generate data given latent variable  $\rightarrow$  **Decoder network**

❖ VAE tries to find the lower bound of  $\log P(X)$  (ELBO (Evidence Lower Bound) )

Objective: **Maximizing**  $\log P(X/z)$  and **minimizing** KL distance between our simple distribution  $Q(z/X)$  and the true latent distribution  $P(z)$ .



# Variational Autoencoder

❖ How to solve?

$$\log P(X) - D_{KL}[Q(z|X) \| P(z|X)] = \underbrace{E[\log P(X|z)]}_{\text{Maximizing } \log P(X/z)} - D_{KL}[Q(z|X) \| P(z)]$$

❖ **Maximizing  $\log P(X/z)$ :** Any discriminative supervised-parametric model (SVM/MLP/...)

❖ Any classifier with  $z$  as input and  $X$  as output, and a proper loss





# Variational Autoencoder

❖ How to solve?

$$\log P(X) - D_{KL}[Q(z|X) \| P(z|X)] = E[\log P(X|z)] - D_{KL}[Q(z|X) \| P(z)]$$

❖ minimizing  $D_{KL}[Q(z|X) \| P(z)]$ :

- $P(z)$ : Latent variable pdf, we will sample it to generate  $X$ .  $\rightarrow P(z) \sim N(0,1)$
- Easy to make  $Q(z|X)$  as-close-as possible to it
- $Q(z|X)$ : Another simplification:  $N(\mu(X), \Sigma(X))$ , and diagonal,  $\Sigma(X)$ !
- Easy to compute KL between two Gaussian!



# Variational Autoencoder

❖ How to solve?

$$\log P(X) - D_{KL}[Q(z|X) \| P(z|X)] = E[\log P(X|z)] - D_{KL}[Q(z|X) \| P(z)]$$

❖ minimizing  $D_{KL}[Q(z|X) \| P(z)]$ :

- $P(z) \sim N(0,1)$ ,  $Q(z|X) \sim N(\mu(X), \Sigma(X))$

$$D_{KL}[N(\mu(X), \Sigma(X)) \| N(0, 1)] = \frac{1}{2} (\text{tr}(\Sigma(X)) + \mu(X)^T \mu(X) - k - \log \det(\Sigma(X)))$$



# Variational Autoencoder

❖ minimizing  $D_{KL}[Q(z|X)||P(z)]$ :

- $P(z) \sim N(0,1)$ ,  $Q(z|X) \sim N(\mu(X), \Sigma(X))$

$$D_{KL}[N(\mu(X), \Sigma(X))||N(0, 1)] = \frac{1}{2} (\text{tr}(\Sigma(X)) + \mu(X)^T \mu(X) - k - \log \det(\Sigma(X)))$$

$$D_{KL}[N(\mu(X), \Sigma(X))||N(0, 1)] = \frac{1}{2} \left( \sum_k \Sigma(X) + \sum_k \mu^2(X) - \sum_k 1 - \log \prod_k \Sigma(X) \right)$$

$$= \frac{1}{2} \left( \sum_k \Sigma(X) + \sum_k \mu^2(X) - \sum_k 1 - \sum_k \log \Sigma(X) \right)$$

$$= \frac{1}{2} \sum_k (\Sigma(X) + \mu^2(X) - 1 - \log \Sigma(X))$$

# Variational Autoencoder

❖ minimizing  $D_{KL}[Q(z|X)\|P(z)]$ :

- $P(z) \sim N(0,1)$ ,  $Q(z|X) \sim N(\mu(X), \Sigma(X))$

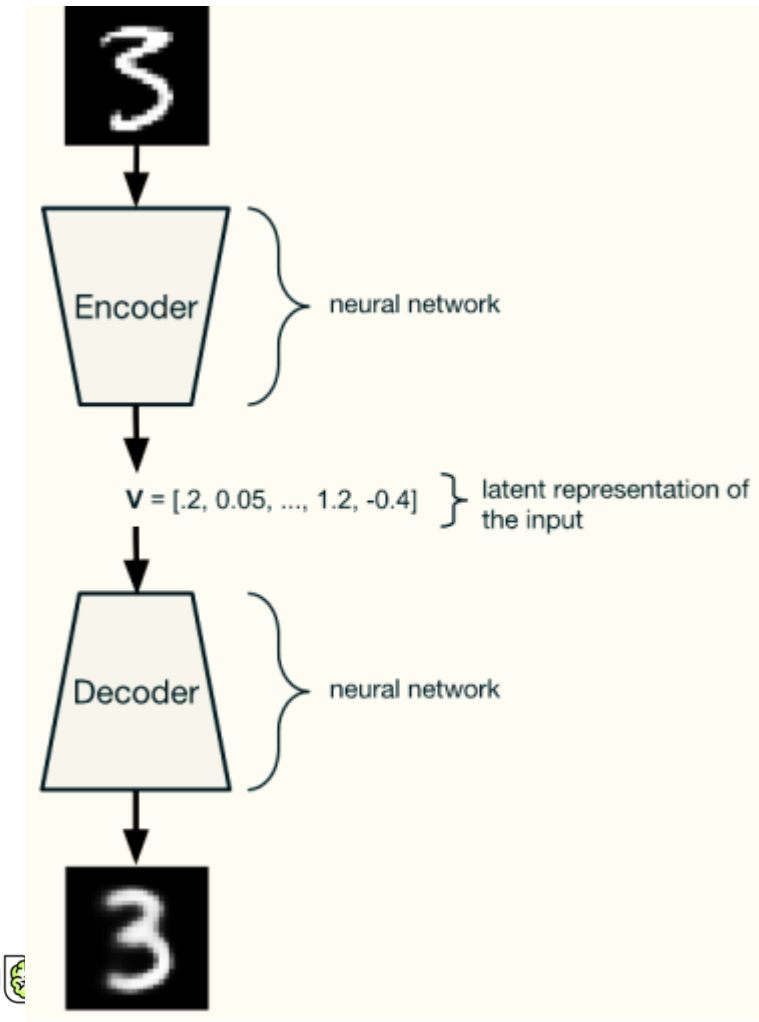
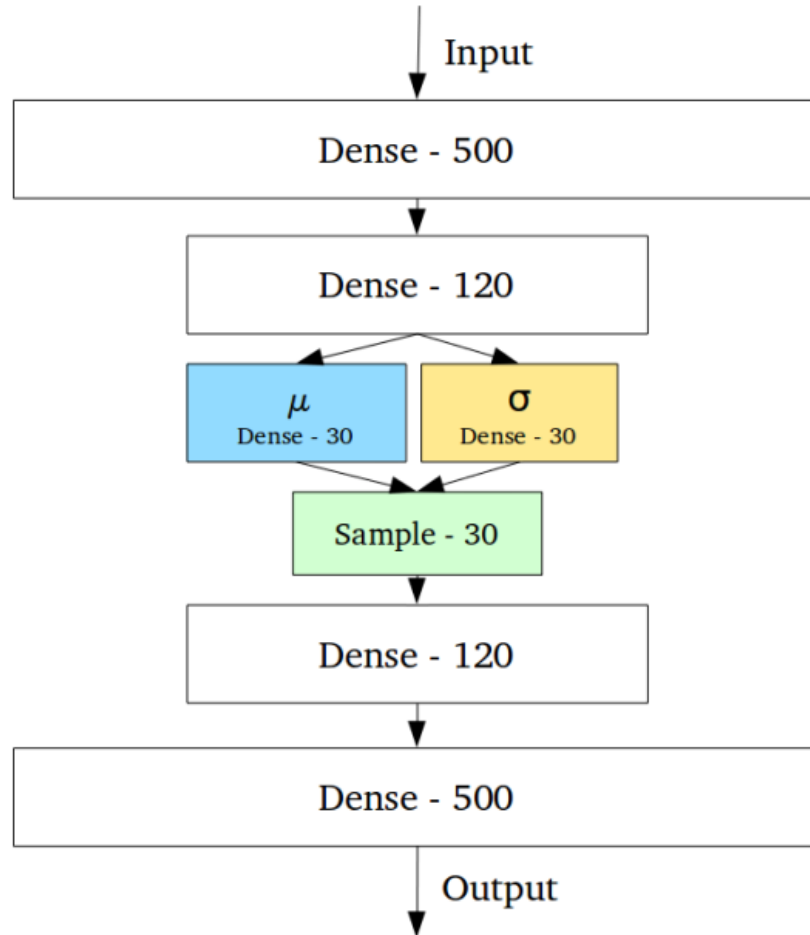
❖ In practice, however, it's better to model  $\Sigma(X)$  as  $\log \Sigma(X)$ , as it is more numerically stable to take exponent compared to computing log.

$$D_{KL}[N(\mu(X), \Sigma(X))\|N(0, 1)] = \frac{1}{2} \sum_k (\exp(\Sigma(X)) + \mu^2(X) - 1 - \Sigma(X))$$




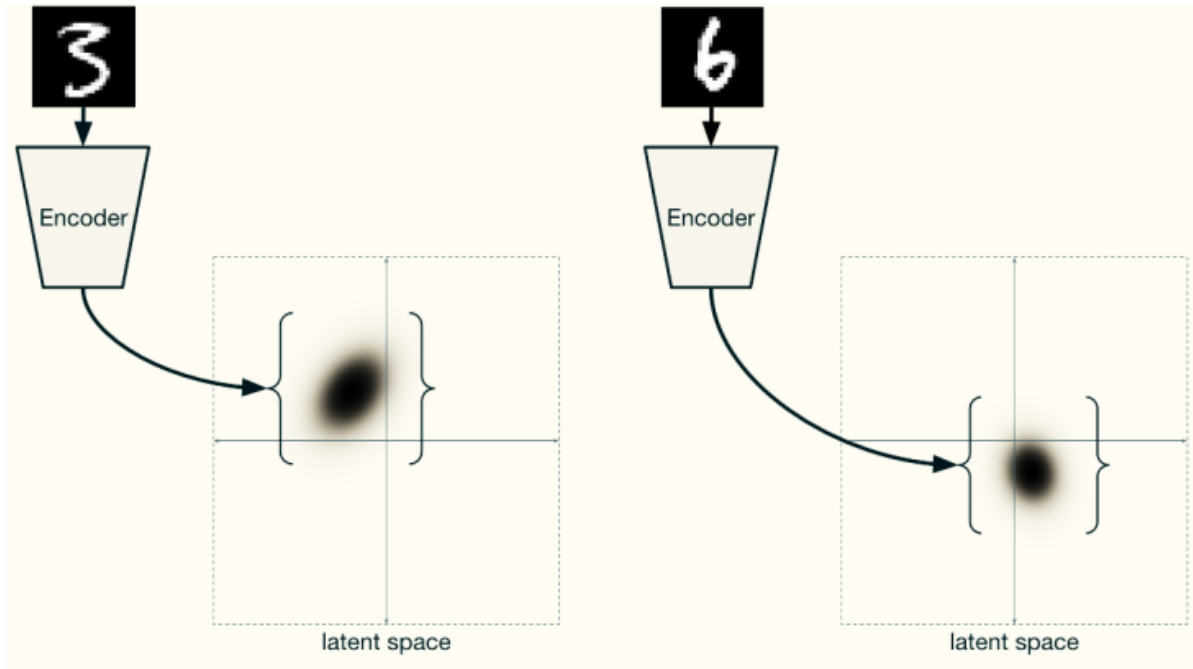
# Variational Autoencoder

## ❖ Overall description



# Variational Autoencoder

## ❖ Overall description



# Variational Autoencoder

## ❖ References

- Auto-encoding variational bayes ([https://arxiv.org/pdf/1312.6114.pdf?source=post\\_page-----](https://arxiv.org/pdf/1312.6114.pdf?source=post_page-----))
- <https://agustinus.kristia.de/techblog/2016/12/10/variational-autoencoder/>
- Attribute2image: Conditional image generation from visual attributes





# Generative adversarial Networks (GANs)




# Generative adversarial Networks

- ❖ Explicitly working with the distribution
  - Variational autoencoders
- ❖ Implicit generative model (Generative Adversarial Networks)
  - trained without even needing to explicitly define a density functions




# Generative adversarial Networks

- ❖ Sample from a simple distribution, e.g. random noise
- ❖ Then, learn transformation to training distribution



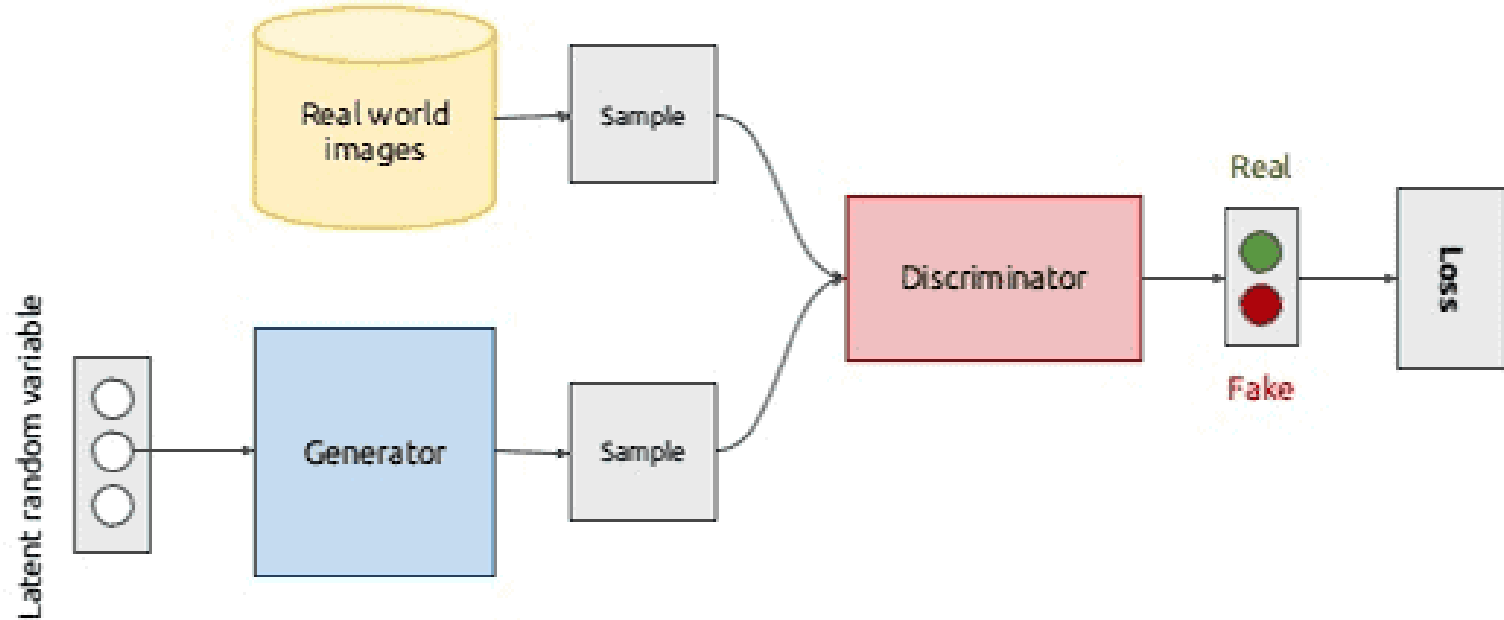
# Generative adversarial Networks

- ❖ Generator network: try to fool the discriminator by generating real-looking images
- ❖ Discriminator network: try to distinguish between real and fake images



# Generative adversarial Networks

- ❖ Generator network: try to fool the discriminator by generating real-looking images
- ❖ Discriminator network: try to distinguish between real and fake images



# Generative adversarial Networks

## ❖ GAN:

- Generative: Learn a Generative Model
- Adversarial: Trained in an Adversarial Setting
- Network: Use Deep Neural Networks

## ❖ GAN final goal:



- Generation of samples from some distribution
- Create Art
- Image-to-Image translation (Aerial images to Maps)





# GANs output examples

Phillip Isola et al., Image-to-Image Translation with Conditional Adversarial Networks, CVPR 2017.





# GANs output examples

C. Ledig et al, Photo-Realistic Single Image Super-Resolution Using a Generative Adversarial Network, CVPR 2017.

original



bicubic  
(21.59dB/0.6423)



SRResNet  
(23.44dB/0.7777)



SRGAN  
(20.34dB/0.6562)



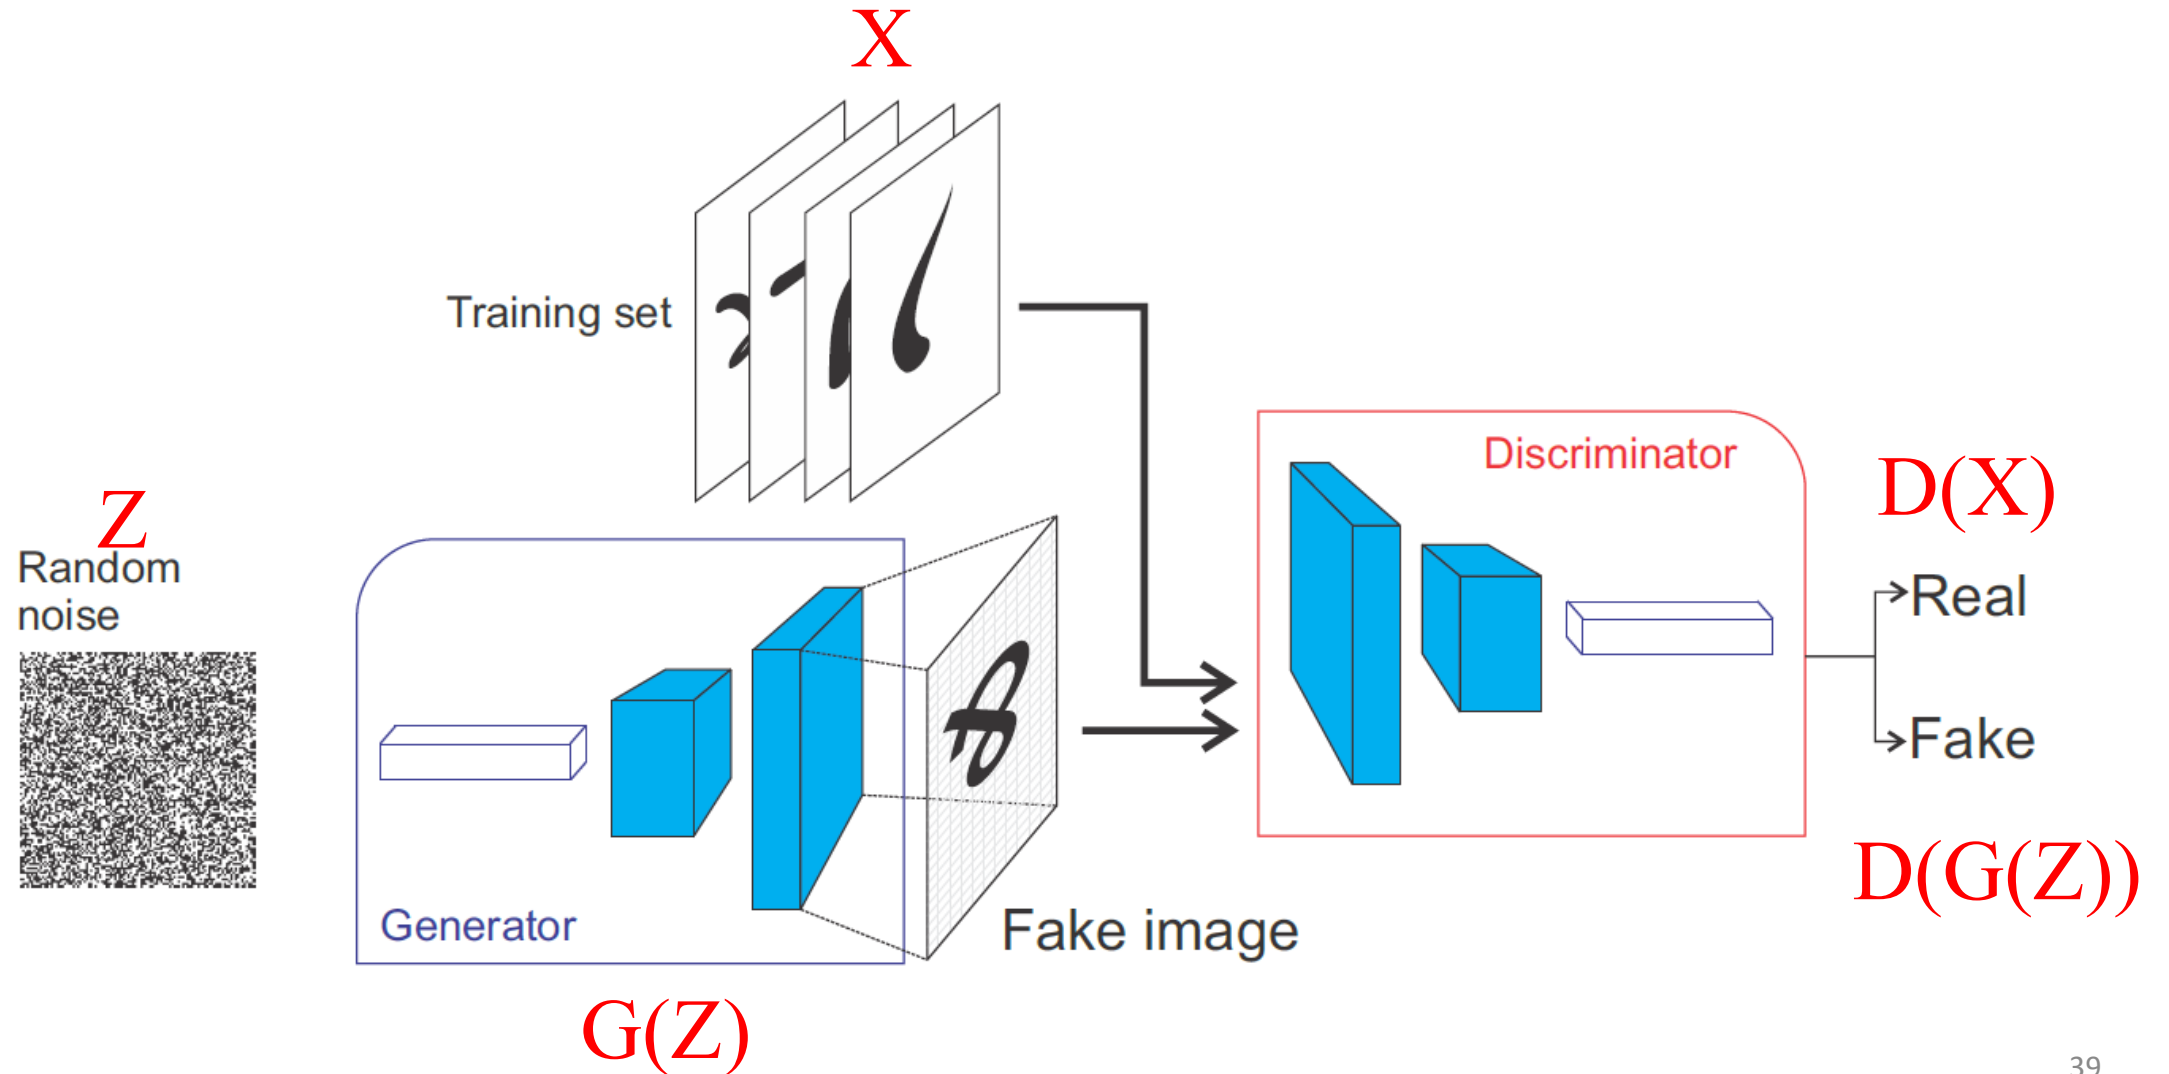
# GANs output examples

<https://thispersondoesnotexist.com/>





# Generative adversarial Networks



# Generative adversarial Networks

- ❖ Generator network: intend to generate real-looking samples
- ❖ Discriminator network: intend to distinguish between real and fake samples

Discriminator output  
for real samples

Discriminator output  
for generated samples

$$J^{(D)} = \mathbb{E}_{x \sim p_{data}} [\log D(x)] + \mathbb{E}_{z \sim p(z)} [\log(1 - D(G(z)))]$$

$$J^{(G)} = \mathbb{E}_{z \sim p(z)} [\log D(G(z))]$$

$$\theta_D^* = \max_{\theta_D} \mathbb{E}_{x \sim p_{data}} [\log D_{\theta_D}(x)] + \mathbb{E}_{z \sim p(z)} [\log(1 - D_{\theta_D}(G_{\theta_G}(z)))]$$

$$\theta_G^* = \max_{\theta_G} \mathbb{E}_{z \sim p(z)} [\log(D_{\theta_D}(G_{\theta_G}(z)))]$$



# Generative adversarial Networks

- ❖ A minimax game between generator and discriminator:

$$\min_{\theta_g} \max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

- ❖ Solving problem:

- ❖ Gradient **ascent** on discriminator

$$\max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

- ❖ Gradient **descent** on generator

$$\min_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))$$



# Generative adversarial Networks

- ❖ A minimax game between generator and discriminator:

$$\min_{\theta_g} \max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

- ❖ Solving problem:

- ❖ Gradient **ascent** on discriminator

$$\max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

- ❖ Gradient **ascent** on generator

$$\max_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(D_{\theta_d}(G_{\theta_g}(z)))$$



# Generative adversarial Networks

## ❖ Training GANs:

---

**Algorithm 1** Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator,  $k$ , is a hyperparameter. We used  $k = 1$ , the least expensive option, in our experiments.

---

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

**end for**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

**end for**

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

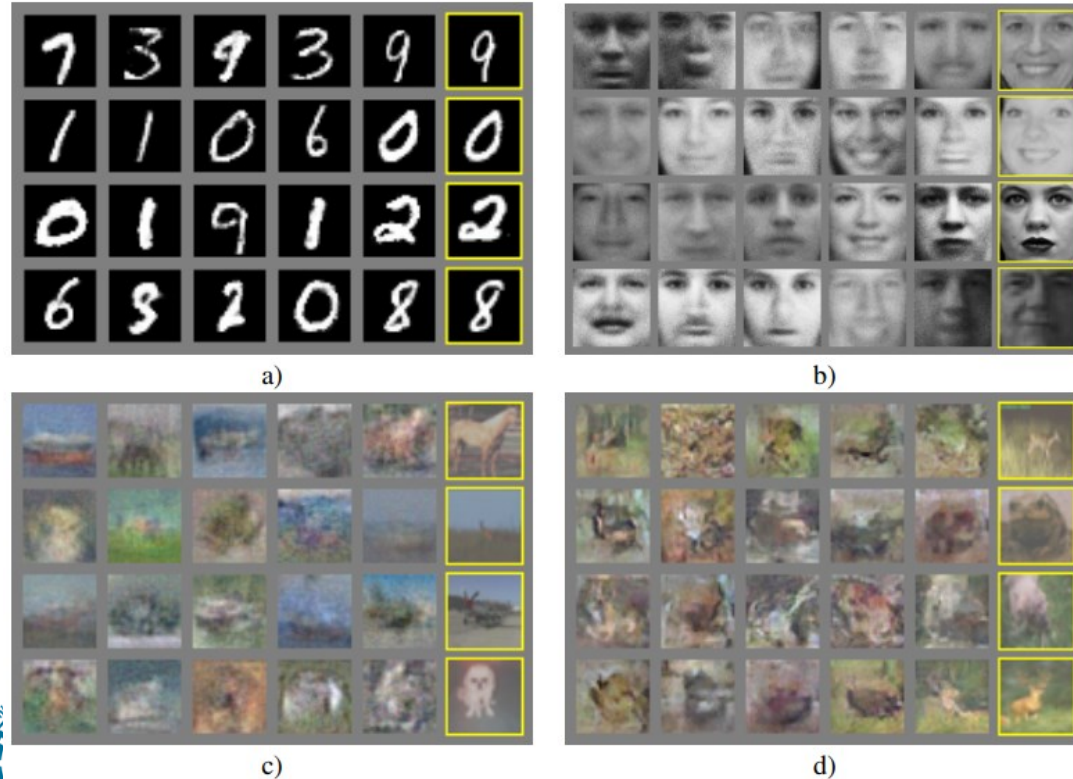
---



# Generative adversarial Networks

## ❖ Generated samples

(Rightmost column shows the nearest training example of the neighboring sample)



# Any Question?

❖ Contact me!!

[Rezakarimzadeh1996@gmail.com](mailto:Rezakarimzadeh1996@gmail.com)





اولین کنگره بین المللی

فناوریهای پیشرفته در حوزه سلامت و کاربردهای هوش مصنوعی در پزشکی

The First International Congress on  
Advanced Health Technologies-Artificial Intelligence in Medicine

